

ENOSS - Event Notifications in OpenStack Swift

Nemanja Vasiljević*



Abstract

Currently object storage OpenStack Swift does not provide any informations to users about events occurred in storage they own/have access to. Users do not have information when content of their object storage is accessed, changed, created or deleted. The aim of this paper is to create solution that will send notifications about events occurred in OpenStack Swift to user specified destinations. Proposed solution, using metadata, allows users to specify where and which event should be published based on even types (read, create, modify, delete) and other properties such as object prefix, suffix, size, etc. It also offers mutiple destinations(Beanstalkd queue, Kafka, etc.) to which notifications can be publshied. Solution is fully compatible with AWS S3 Event Notifications and compared to AWS supports more destinations, event types, filters and allows unsuccessful events to be published as well. Event notification can be used not just for monitoring, but also for automatization and serverless computing similarly to AWS Lambda.

Keywords: Event — Notifications — OpenStack Swift

Supplementary Material: [Demonstration Video](#) — [Downloadable Code](#)

*xvasil03@stud.fit.vutbr.cz, Faculty of Information Technology, Brno University of Technology

1. Introduction

Object storage is data storage architecture that manages data as objects, and each object typically includes data itself and some additional information stored in objects metadata. Since object storages are often used in cloud computing, data are stored in remote locations, to which users do not have direct and complete access, some users or external services might want to receive information about specific events in storages where their data are located. For example, currently there is no easy way to detect changes in specific container, except to list it's content and compare timestamps, which can be complex, slow and unefficient if there is a lot of objects in storage.

The importance of this work is to provide event information to users in OpenStack Swift, which will

allow users to react to those events, create more sophisticated backend operations, postprocessing and automatization, or possibly prevent/detect unwanted actions. In addition, providing event notifications will allow users to have a better picture of what is going on in their storage and improve monitoring in object storage.

Users can be interested in only specific events, for example creating of new objects in container, therefore it is important that proposed solution allows event filtering based on event type and other properties (e.g. object name prefix/suffix/size). Since object storage have multiple users, each user can have different requirements for event notification and proposed solution must be prepared for it.

Application of event notifications is various, from

17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32

33 simple monitoring or webhook, to more sophisticated
34 application as serverless computing like AWS Lambda.
35 This means that structure of event notification may dif-
36 fer based on its application and destination to which is
37 published. Proposed solution must be ready to publih
38 event notification to different destinations as well as in
39 different event notification structure.

40 AWS S3 object storage is one of the most popu-
41 lar storage with their own API, that is supported by
42 many other object storages including OpenStack Swift.
43 Since AWS S3 supports event notifications, it would be
44 ideal if proposed solution in OpenStack is compatible
45 with S3 event notification protocol. This would allow
46 easier transfer users from AWS S3 to OpenStack Swift.
47 As result, not only that OpenStack Swift would offer
48 same functionality as AWS S3 (that currently lacks),
49 but protocol would be compatible with AWS S3, which
50 would allow easier tranfer users from AWS S3 to Open-
51 Stack Swift. Therefore users would not have to learn
52 additional protocol, and can follow existing AWS S3
53 which is most popular and well documented.

54 2. OpenStack Swift

55 OpenStack Swift is open-source object storage devel-
56 oped by Rackspace, a company that, together with
57 NASA, created the OpenStack project. After becom-
58 ing an open-source project, Swift became the leading
59 open-source object storage supported and developed
60 by many famous IT companies, such as Red Hat, HP,
61 Intel, IBM, and others.

62 OpenStack Swift is a multi-tenant, scalable, and
63 durable object storage capable of storing large amounts
64 of unstructured data at low cost[?].

65 2.1 Data model

66 OpenStack Swift allows users to store unstructured
67 data objects with a canonical name containing *account*,
68 *container* and *object* in given order[?]. The account
69 names must be unique in the cluster, the container
70 name must be unique in the account space, and the ob-
71 ject names must be unique in the container. Other than
72 that, if containers have the same name but belong to a
73 different account, then they represent different storage
74 locations. The same principle applies to objects. If
75 objects have the same name but not the same container
76 and account name, then these objects are different.

77 **Accounts** are root storage locations for data. Each
78 account contains a list of containers within the account
79 and metadata stored as key-value pairs. Accounts are
80 stored in the account database. In OpenStack Swift, ac-
81 count is **storage account** (more like storage location)
82 and **do not represent a user identity**[?].

Containers are user-defined storage locations in 83
the account namespace where objects are stored. Con- 84
tainers are one level below accounts, therefore they 85
are not unique in the cluster. Each container has a list 86
of objects within the container and metadata stored 87
as key-value pairs. Containers are stored in container 88
database[?]. 89

Objects represent data stored in OpenStack Swift. 90
Each object belongs to one (and only one) container. 91
An object can have metadata stored as key-value pairs. 92
Swift stores multiple copies of an object across the 93
cluster to ensure durability and availability. Swift 94
does this by assigning an object to *partition*, which is 95
mapped to multiple drives, and each driver will contain 96
object copy[?]. 97

52 Main processes

The path towards data in OpenStack Swift consists of 99
four main software services: **Proxy server**, **Account** 100
server, **Continaer server** and **Object server**. Typi- 101
cally Account, Container and Object server are located 102
on same machine creating **Storage node**. 103

Proxy server is the service responsible for com- 104
munication with external clients. For each request, it 105
will look up storage location(node) for an account, con- 106
tainer, or object and route the request accordingly[?]. 107
The proxy server is responsible for handling many fail- 108
ures. For example, when a client sends a PUT request 109
to OpenStack Swift, the proxy server will determine 110
which nodes store the object. If some node fails, a 111
proxy server will choose a hand-off node to write data. 112
When a majority of nodes respond successfully, then 113
the server proxy will return a success response code[?]. 114

Account server stores information about contain- 115
ers in a particular account to SQL database. It is 116
responsible for listing containers. It does not know 117
where specific containers are, just what containers are 118
in an account[?]. 119

Container server is similar to account server, ex- 120
cept it is responsible for listing objects and also does 121
not know where specific objects are[?]. 122

Object Server is blob storage capable of storing, 123
retrieving, and deleting objects. Objects are stored 124
as binary files to a filesystem, where metadata are 125
stored in the *file's extended attributes (xattrs)*. This 126
requires a filesystem with support of such attributes. 127
Each object is stored using a hash value of object path 128
(account/container/object) and timestamp. This allows 129
storing multiple versions of an object. Since last write 130
wins (due to timestamp), it is ensured that the correct 131
object version is served[?]. 132




Figure 1. OpenStack Swift servers architecture.

and call the next middleware in the pipeline or intercept a request. In the healthcheck example, if the path directs to `/healthcheck`, the middleware will return HTTP Response, and other middlewares in the pipeline will not be called.

Method `filter_factory` is used by the Python Paste framework to instantiate middleware.

```

import os
from swift.common.swob import Request, Response

class HealthCheckMiddleware(object):
    def __init__(self, app, conf):
        self.app = app

    def __call__(self, env, start_response):
        req = Request(env)
        if req.path == '/healthcheck':
            return Response(request=req, body=b"OK", content_type="text/plain")(env, start_response)
        return self.app(env, start_response)

def filter_factory(global_conf, **local_conf):
    local_conf = local_conf
    conf = global_conf.copy()
    conf.update(local_conf)

    def healthcheck_filter(app):
        return HealthCheckMiddleware(app, conf)
    return healthcheck_filter

```

Listing 2. Example of healthcheck middleware in OpenStack Swift

2.4 Metadata

OpenStack Swift separates metadata into 3 categories based on their use:

User Metadata - User metadata takes form

`X-<type>-Meta-<key>:<value>`

where `<type>` represent resource type(i.e. account, container, object), and `<key>` and `<value>` are set by user. User metadata remain persistent until are updated using new value or removed using header `X-<type>-Meta-<key>` with no value or a header `X-Remove-<type>-Meta-<key>:<ignored-value>`.

System Metadata - System metadata takes form `X-<type>-Sysmeta-<key>:<value>` where `<type>` represent resource type(i.e. account, container, object) and `<key>` and `<value>` are set by internal service in Swift WSGI Server. All headers containing system metadata are deleted from a client request. System metadata are visible only inside Swift, providing a means to store potentially sensitive information regarding Swift resources.

Object Transient-Sysmeta - This type of

2.3 Middleware

Using Python WSGI middleware, users can add functionalities and behaviors to OpenStack Swift. Most middlewares are added to the Proxy server but can also be part of other servers (account server, container server, or object server).

Middleware are added by changing the configuration of servers. Listing 1 shows how to add *webhook middleware* to proxy server by changing its pipeline (*pipeline:main*). Middlewares are executed in the given order (first will be called *webhook middleware*, then *proxy-server middleware*).

Some of the middlewares are required and will be automatically inserted by swift code[?].

Listing 1. Example of proxy server configuration (proxy-server.conf).

```

[DEFAULT]
log_level = DEBUG
user = <your-user-name>

[pipeline:main]
pipeline = webhook proxy-server

[filter:webhook]
use = egg:swift#webhook

[app:proxy-server]
use = egg:swift#proxy

```

Interface - OpenStack Swift servers are implemented using Python WSGI applications. Therefore only Python WSGI middlewares are accepted in OpenStack Swift.

Listing 2 provides example of simplified *healthcheck middleware*. The constructor takes two arguments, the first is a WSGI application, and the second is a configuration of middleware defined using Python Paste framework in *proxy-server.conf*. Middleware must have a call method containing the request environment information and response from previously called middleware. Middleware can perform some operations

224 metadata have form of X-Object-Transient-
 225 -Sysmeta-<key>:<value>. Transient-sysmeta
 226 has a similar behavior as system metadata and can
 227 be accessed only within Swift, and headers contain-
 228 ing Transient-sysmeta are dropped. If middleware
 229 wants to store object metadata, it should use transient-
 230 sysmeta[?].

231 3. Existing solutions

232 There is no official OpenStack solution that satisfies all
 233 requirements mentioned in section 1, although some of
 234 existing programs can be used to partially solve some
 235 of the problems.

236 **Webhook middleware** described in ?? can be
 237 used for detection of new objects in specific container.
 238 With some tweaks it could be able to detect object
 239 deletion and modifying too. One of many limitations
 240 of this middleware is lack of support different des-
 241 tinations (it can publish notification only to single
 242 type of destination), no filtering, single type of event
 243 notification structure and incompatibility with AWS
 244 S3.

245 **OpenStack Swift attempts** - OpenStack Swift is
 246 aware of lack of event notifications and in order to
 247 solve it they created specification for this problem[?].
 248 This specification was mainly focused on detection
 249 changes inside specific container (creation, modifying
 250 and deletion of objects). There were two attempts to
 251 solve this problem.

- 252 • **First attempt??** - allowed sending notifications
 253 only to Zahar[TODO] queue and had very simple
 254 event notification structure. Notification
 255 contained only informations about names of ac-
 256 count, container and object on which event oc-
 257 curred and name of HTTP method.
- 258 • **Second attempt??** - was more sophisticated so-
 259 lution that was designed to support multiple des-
 260 tinations to which notification can be published.
 261 Event notification structure was expanded for in-
 262 formations such as eTag (MD5 checksum) and
 263 transaction id. Author introduced concept of
 264 "notification policy", which represented config-
 265 uration of event notifications. One of main cri-
 266 tiques made by code reviewers was incompati-
 267 bility with AWS S3 storage.

268 Both attempts are outdated and due to lack of in-
 269 terest from users/operators OpenStack Swift halted
 270 development for this problem.

271 **ENOSS** - my solution, code name ENOSS, satis-
 272 fies all requirements specified in section 1. Key fea-
 273 tures are: events filtering, support of multiple desti-

nations, AWS S3 compatibility, different event notifi- 274
 cation structure, definition of interfaces for future ex- 275
 pansions of filters, destinations and event notification 276
 structure, and design that allows its effortless expan- 277
 sions. 278

4. ENOSS 279

ENOSS (Event Notifications in OpenStack Swift) is 280
 program that enables publishing notifications contain- 281
 ing information about occurred events in OpenStack 282
 Swift. It is implemented in form of Python WSGI mid- 283
 dleware and is located in Proxy server pipeline. Since 284
 Proxy server communicates with external users, by 285
 placing ENOSS in its pipeline ENOSS is able to react 286
 on every request made by users to OpenStack Swift, 287
 which makes Proxy server ideal places for ENOSS. 288

4.1 Key features 289

The proposed middleware heavily utilizes container- 290
 s/buckets and accounts metadata. Information spec- 291
 ifying which event should be published and where 292
 is stored in metadata of upper level. For publishing 293
 events regarding objects, the configuration is stored 294
 in a container metadata and for container events, the 295
 configuration is stored in an account level. 296

Multi user environment - since many different 297
 users communicate with OpenStack Swift each of 298
 them can be interested in different event notifications. 299
 ENOSS solves this problem by allowing each container 300
 and account to have their own notification configura- 301
 tion. 302

Event filtering - one of the main requirements 303
 for event notifications is allowing users to specify 304
 for which events should notifications be published - i.e. 305
 event filtering. ENOSS allows user to specify which 306
 types of events should be published (object/container 307
 creation, deletion, access, etc.). ENOSS goes little 308
 bit further and also allows users to specify rules that 309
 must be satisfied in order for event notification to be 310
 published. Some of the rule operators are: object/con- 311
 tainer name prefix/suffix and object size. Using this 312
 feature user can select for example, only events regard- 313
 ing objects bigger than 50mb (operator: object size) or 314
 events regarding pictures (operator: object suffix). 315

Multiple destinations - since event notifications 316
 has multiple applications, from monitoring to automa- 317
 tization, it is important that proposed solution is able 318
 to publish notification to multiple different destinations. 319
 ENOSS is fully capable of publishing event notifica- 320
 tions to many different destinations (e.g. Beanstalkd 321
 queue, Kafka). In ENOSS, publishing notification 322
 about single event is not limited to only one destina- 323

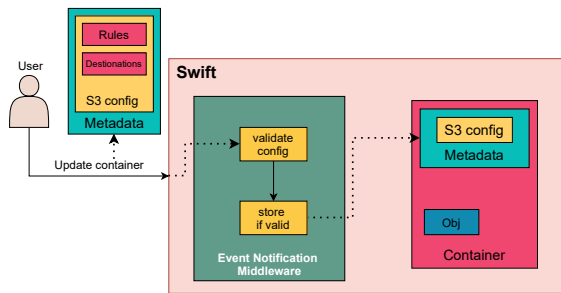


Figure 2. Process of setting event notification configuration in ENOSS.

tion, but, if user wishes, it can be published to multiple different destination per single event. This feature allows event notification to be used for multiple applications simultaneously.

Event notification structure - depending on application of event notification structure of notification may differ. ENOSS supports several different notification structures and using event notification configuration users are able to select type of event notification structure.

AWS S3 compatibility - ENOSS puts big emphasis on support and compatibility with AWS S3. Structure of event configuration and event names in ENOSS are compatible with AWS S3. ENOSS also supports all filtering rules from AWS S3 and default event notification structure is compatible with AWS S3. This is all done in order to ease transfer users from AWS S3 to OpenStack Swift and using existing protocol, that is well documented, users will have easier time learning and using event notifications in OpenStack Swift.

4.2 Configuration

Setting event notification configuration - in order to enable event notifications on specific container, first step is to store its configuration. For this purpose ENOSS uses API:

POST /v1/<acc>/<cont>?notification

Figure 2 describes process of storing event configuration. Authorized user sends event notification configuration using request body, ENOSS perform validation, if configuration is valid, ENOSS will store configuration to container system metadata, otherwise it will return unsuccessful HTTP code.

Reading stored event notification configuration

- Event notification configuration can contain sensitive information. Since ENOSS uses stores configuration to storage using system metadata, which can be access only by application within OpenStack Swift, it disable reading stored configuration by simple GET/HEAD requests. For this purpose ENOSS offer API

GET /v1/<acc>/<cont>?notification

For security reasons, ENOSS allow only users with

write rights to read stored configuration.

Configuration structure - Listing 3 describes event notification configuration. <Target> represent targeted destination where event notifications will be sent (e.g., Beanstalkd, Elasticsearch). <FilterKey> is a unique name of a filter containing rules that must be satisfied in order to publish events.

Event type takes form :

s3:<Type><Action>:<Method>

and are compatible with Amazon S3 event types. Type represents resource type (object, bucket), action represent action performed by user and can have values: Created, Removed, Accessed. The method represents the REST API method performed by a user: Get, Put, Post, Delete, Copy, Head. For example, if a new object was created, even type would be described as s3:ObjectCreated:Put. To match event type regardless of API method assign value * to <Method>.

Listing 3. Structure of event notification configuration

```
{
  "<Target>Configurations": [
    {
      "Id": "configuration id",
      "TargetParams": "set of key-value
        pairs, used specify dynamic
        parameters of targeted
        destination (e.g., name of
        beanstalkd tube or name of the
        index in Elasticsearch)",
      "Events": "array of event types that
        will be published",
      "PayloadStructure": "type of event
        notification structure: S3 or
        CloudEvents (default value S3)",
      "Filter": {
        "<FilterKey>": {
          "FilterRules": [
            {
              "Name": "filter operations (i
                .e. prefix, suffix, size)
              ",
              "Value": "filter value"
            }
          ]
        }
      }
    }
  ]
}
```

4.3 Interfaces

Destination -

Payload -

Filter Rule -

Supported event types

421 4.4 Performance

422 5. Conclusions

423 **[Paper Summary]** What was the paper about, then?
424 What the reader needs to remember about it? Lorem
425 ipsum dolor sit amet, consectetur adipiscing elit. Proin
426 vitae aliquet metus. Sed pharetra vehicula sem ut var-
427 ius. Aliquam molestie nulla et mauris suscipit, ut
428 commodo nunc mollis.

429 **[Highlights of Results]** Exact numbers. Remind
430 the reader that the paper matters. Lorem ipsum do-
431 lor sit amet, consectetur adipiscing elit. Sed tempus
432 fermentum ipsum at venenatis. Curabitur ultricies,
433 mauris eu ullamcorper mattis, ligula purus dapibus mi,
434 vel dapibus odio nulla et ex. Sed viverra cursus mattis.
435 Suspendisse ornare semper condimentum. Interdum et
436 malesuada fames ac ante ipsum.

437 **[Paper Contributions]** What is the original con-
438 tribution of this work? Two or three thoughts that one
439 should definitely take home. Lorem ipsum dolor sit
440 amet, consectetur adipiscing elit. Praesent posuere
441 mattis ante at imperdiet. Cras id tincidunt purus. Ali-
442 quam erat volutpat. Morbi non gravida nisi, non iaculis
443 tortor. Quisque at fringilla neque.

444 **[Future Work]** How can other researchers / devel-
445 opers make use of the results of this work? Do you
446 have further plans with this work? Or anybody else?
447 Lorem ipsum dolor sit amet, consectetur adipiscing
448 elit. Suspendisse sollicitudin posuere massa, non con-
449 vallis purus ultricies sit amet. Duis at nisl tincidunt,
450 maximus risus a, aliquet massa. Vestibulum libero
451 odio, condimentum ut ex non, eleifend.

452 Acknowledgements

453 I would like to thank my supervisor X. Y. for her help.