

# ENOSS - Event Notifications in OpenStack Swift

Nemanja Vasiljević\*



## Abstract

Currently object storage OpenStack Swift does not provide any informations to users about events occurred in storage they own/have access to. Users do not have information when content of their object storage is accessed, changed, created or deleted. The aim of this paper is to create solution that will send notifications about events occurred in OpenStack Swift to user specified destinations. Proposed solution, using metadata, allows users to specify where and which event should be published based on even types (read, create, modify, delete) and other properties such as object prefix, suffix, size, etc. It also offers mutiple destinations(Beanstalkd queue, Kafka, etc.) to which notifications can be publshied. Solution is fully compatible with AWS S3 Event Notifications and compared to AWS supports more destinations, event types, filters and allows unsuccessful events to be published as well. Event notification can be used not just for monitoring, but also for automatization and serverless computing similarly to AWS Lambda.

**Keywords:** Event — Notifications — OpenStack Swift

**Supplementary Material:** [Demonstration Video](#) — [Downloadable Code](#)

\*[xvasil03@stud.fit.vutbr.cz](mailto:xvasil03@stud.fit.vutbr.cz), Faculty of Information Technology, Brno University of Technology

## 1. Introduction

Object storage is data storage architecture that manages data as objects, and each object typically includes data itself and some additional information stored in objects metadata. Since object storages are often used in cloud computing, data are stored in remote locations, to which users do not have direct and complete access, some users or external services might want to receive information about specific events in storages where their data are located. For example, currently there is no easy way to detect changes in specific container, except to list it's content and compare timestamps, which can be complex, slow and unefficient if there is a lot of objects in storage.

The importance of this work is to provide event information to users in OpenStack Swift, which will

allow users to react to those events, create more sophisticated backend operations, postprocessing and automatization, or possibly prevent/detect unwanted actions. In addition, providing event notifications will allow users to have a better picture of what is going on in their storage and improve monitoring in object storage.

Users can be interested in only specific events, for example creating of new objects in container, therefore it is important that proposed solution allows event filtering based on event type and other properties (e.g. object name prefix/suffix/size). Since object storage have multiple users, each user can have different requirements for event notification and proposed solution must be prepared for it.

Application of event notifications is various, from

17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32

33 simple monitoring or webhook, to more sophisticated  
34 application as serverless computing like AWS Lambda.  
35 This means that structure of event notification may dif-  
36 fer based on its application and destination to which is  
37 published. Proposed solution must be ready to publish  
38 event notification to different destinations as well as in  
39 different event notification structure.

40 AWS S3 object storage is one of the most popular  
41 storage with their own API, that is supported by  
42 many other object storages including OpenStack Swift.  
43 Since AWS S3 supports event notifications, it would be  
44 ideal if proposed solution in OpenStack is compatible  
45 with S3 event notification protocol. This would allow  
46 easier transfer users from AWS S3 to OpenStack Swift.  
47 As result, not only that OpenStack Swift would offer  
48 same functionality as AWS S3 (that currently lacks),  
49 but protocol would be compatible with AWS S3, which  
50 would allow easier transfer users from AWS S3 to Open-  
51 Stack Swift. Therefore users would not have to learn  
52 additional protocol, and can follow existing AWS S3  
53 which is most popular and well documented.

## 54 2. OpenStack Swift

55 OpenStack Swift is open-source object storage devel-  
56 oped by Rackspace, a company that, together with  
57 NASA, created the OpenStack project. After becom-  
58 ing an open-source project, Swift became the leading  
59 open-source object storage supported and developed  
60 by many famous IT companies, such as Red Hat, HP,  
61 Intel, IBM, and others.

62 OpenStack Swift is a multi-tenant, scalable, and  
63 durable object storage capable of storing large amounts  
64 of unstructured data at low cost[?].

### 65 2.1 Data model

66 OpenStack Swift allows users to store unstructured  
67 data objects with a canonical name containing *account*,  
68 *container* and *object* in given order[?]. The account  
69 names must be unique in the cluster, the container  
70 name must be unique in the account space, and the ob-  
71 ject names must be unique in the container. Other than  
72 that, if containers have the same name but belong to a  
73 different account, then they represent different storage  
74 locations. The same principle applies to objects. If  
75 objects have the same name but not the same container  
76 and account name, then these objects are different.

77 **Accounts** are root storage locations for data. Each  
78 account contains a list of containers within the account  
79 and metadata stored as key-value pairs. Accounts are  
80 stored in the account database. In OpenStack Swift, ac-  
81 count is **storage account** (more like storage location)  
82 and **do not represent a user identity**[?].

**Containers** are user-defined storage locations in 83  
the account namespace where objects are stored. Con- 84  
tainers are one level below accounts, therefore they 85  
are not unique in the cluster. Each container has a list 86  
of objects within the container and metadata stored 87  
as key-value pairs. Containers are stored in container 88  
database[?]. 89

**Objects** represent data stored in OpenStack Swift. 90  
Each object belongs to one (and only one) container. 91  
An object can have metadata stored as key-value pairs. 92  
Swift stores multiple copies of an object across the 93  
cluster to ensure durability and availability. Swift 94  
does this by assigning an object to *partition*, which is 95  
mapped to multiple drives, and each driver will contain 96  
object copy[?]. 97

## 52 Main processes


The path towards data in OpenStack Swift consists of 98  
four main software services: **Proxy server**, **Account** 99  
**server**, **Container server** and **Object server**. Typi- 100  
cally Account, Container and Object server are located 101  
on same machine creating **Storage node**. 102  
103

**Proxy server** is the service responsible for com- 104  
munication with external clients. For each request, it 105  
will look up storage location(node) for an account, con- 106  
tainer, or object and route the request accordingly[?]. 107  
The proxy server is responsible for handling many fail- 108  
ures. For example, when a client sends a PUT request 109  
to OpenStack Swift, the proxy server will determine 110  
which nodes store the object. If some node fails, a 111  
proxy server will choose a hand-off node to write data. 112  
When a majority of nodes respond successfully, then 113  
the server proxy will return a success response code[?]. 114

**Account server** stores information about contain- 115  
ers in a particular account to SQL database. It is 116  
responsible for listing containers. It does not know 117  
where specific containers are, just what containers are 118  
in an account[?]. 119

**Container server** is similar to account server, ex- 120  
cept it is responsible for listing objects and also does 121  
not know where specific objects are[?]. 122

**Object Server** is blob storage capable of storing, 123  
retrieving, and deleting objects. Objects are stored 124  
as binary files to a filesystem, where metadata are 125  
stored in the *file's extended attributes (xattrs)*. This 126  
requires a filesystem with support of such attributes. 127  
Each object is stored using a hash value of object path 128  
(account/container/object) and timestamp. This allows 129  
storing multiple versions of an object. Since last write 130  
wins (due to timestamp), it is ensured that the correct 131  
object version is served[?]. 132



**Figure 1.** OpenStack Swift servers architecture.

and call the next middleware in the pipeline or intercept a request. In the healthcheck example, if the path directs to `/healthcheck`, the middleware will return HTTP Response, and other middlewares in the pipeline will not be called.

Method `filter_factory` is used by the Python Paste framework to instantiate middleware.

```

import os
from swift.common.swob import Request, Response

class HealthCheckMiddleware(object):
    def __init__(self, app, conf):
        self.app = app

    def __call__(self, env, start_response):
        req = Request(env)
        if req.path == '/healthcheck':
            return Response(request=req, body=b"OK", content_type="text/plain")(env, start_response)
        return self.app(env, start_response)

def filter_factory(global_conf, **local_conf):
    conf = global_conf.copy()
    conf.update(local_conf)

    def healthcheck_filter(app):
        return HealthCheckMiddleware(app, conf)
    return healthcheck_filter

```

**Listing 2.** Example of healthcheck middleware in OpenStack Swift

## 2.4 Metadata

OpenStack Swift separates metadata into 3 categories based on their use:

**User Metadata** - User metadata takes form

`X-<type>-Meta-<key>:<value>`

where *<type>* represent resource type(i.e. account, container, object), and *<key>* and *<value>* are set by user. User metadata remain persistent until are updated using new value or removed using header `X-<type>-Meta-<key>` with no value or a header `X-Remove-<type>-Meta-<key>:<ignored-value>`.

**System Metadata** - System metadata takes form `X-<type>-Sysmeta-<key>:<value>`

where *<type>* represent resource type(i.e. account, container, object) and *<key>* and *<value>* are set by internal service in Swift WSGI Server. All headers containing system metadata are deleted from a client request.

System metadata are visible only inside Swift, providing a means to store potentially sensitive information regarding Swift resources.

## 2.3 Middleware

Using Python WSGI middleware, users can add functionalities and behaviors to OpenStack Swift. Most middlewares are added to the Proxy server but can also be part of other servers (account server, container server, or object server).

Middleware are added by changing the configuration of servers. Listing 1 shows how to add *webhook middleware* to proxy server by changing its pipeline (*pipeline:main*). Middlewares are executed in the given order (first will be called *webhook middleware*, then *proxy-server middleware*).

Some of the middlewares are required and will be automatically inserted by swift code[?].

**Listing 1.** Example of proxy server configuration (proxy-server.conf).

```

[DEFAULT]
log_level = DEBUG
user = <your-user-name>

[pipeline:main]
pipeline = webhook proxy-server

[filter:webhook]
use = egg:swift#webhook

[app:proxy-server]
use = egg:swift#proxy

```

**Interface** - OpenStack Swift servers are implemented using Python WSGI applications. Therefore only Python WSGI middlewares are accepted in OpenStack Swift.

Listing 2 provides example of simplified *healthcheck middleware*. The constructor takes two arguments, the first is a WSGI application, and the second is a configuration of middleware defined using Python Paste framework in *proxy-server.conf*. Middleware must have a call method containing the request environment information and response from previously called middleware. Middleware can perform some operations

224 **Object Transient-Sysmeta** - This type of  
225 metadata have form of X-Object-Transient-  
226 -Sysmeta-<key>:<value>. Transient-sysmeta  
227 has a similar behavior as system metadata and can  
228 be accessed only within Swift, and headers contain-  
229 ing Transient-sysmeta are dropped. If middleware  
230 wants to store object metadata, it should use transient-  
231 sysmeta[?].

### 232 3. Existing solutions

233 There is no official OpenStack solution that satisfies all  
234 requirements mentioned in section 1, although some of  
235 existing programs can be used to partially solve some  
236 of the problems.

237 **Webhook middleware** described in ?? can be  
238 used for detection of new objects in specific container.  
239 With some tweaks it could be able to detect object  
240 deletion and modifying too. One of many limitations  
241 of this middleware is lack of support different des-  
242 tinations (it can publish notification only to single  
243 type of destination), no filtering, single type of event  
244 notification structure and incompatibility with AWS  
245 S3.

246 **OpenStack Swift attempts** - OpenStack Swift is  
247 aware of lack of event notifications and in order to  
248 solve it they crated specification for this problem[?].  
249 This specification was mainly focused on detection  
250 changes inside specific container (creation, modifying  
251 and deletiton of objects). There were two attempts to  
252 solve this problem.

- 253 • **First attempt??** - allowed sending notifications  
254 only to Zahar[TODO] queue and had very simple  
255 event notification structure. Notification  
256 contained only informations about names of ac-  
257 count, container and object on which event oc-  
258 curred and name of HTTP method.
- 259 • **Second attempt??** - was more sophisticated so-  
260 lution that was design to support multiple des-  
261 tinations to which notfication can be published.  
262 Event notification structure was expanded for in-  
263 formations such as eTag (MD5 checksum) and  
264 transaction id. Author introduced concept of  
265 "notification policy", which represented config-  
266 uration of event notifications. One of main cri-  
267 tiques made by code reviewers was incompati-  
268 bility with AWS S3 storage.

269 Both attempts are outdated and due to lack of in-  
270 terest from users/operators OpenStack Swift halted  
271 development for this problem.

272 **ENOSS** - my solution, code name ENOSS, satis-  
273 fies all requirements specified in section 1. Key fea-

tures are: events filtering, support of multiple desti- 274  
nations, AWS S3 compatibility, different event noti- 275  
fication structure, definition of interfaces for future 276  
expansions and different filters, destinations and event 277  
notification structure, and design that allows its effort- 278  
less expansions. 279

## 4. ENOSS 280

## 5. Conclusions 281

**[Paper Summary]** What was the paper about, then? 282  
What the reader needs to remember about it? Lorem 283  
ipsum dolor sit amet, consectetur adipiscing elit. Proin 284  
vitae aliquet metus. Sed pharetra vehicula sem ut var- 285  
ius. Aliquam molestie nulla et mauris suscipit, ut 286  
commodo nunc mollis. 287

**[Highlights of Results]** Exact numbers. Remind 288  
the reader that the paper matters. Lorem ipsum do- 289  
lor sit amet, consectetur adipiscing elit. Sed tempus 290  
fermentum ipsum at venenatis. Curabitur ultricies, 291  
mauris eu ullamcorper mattis, ligula purus dapibus mi, 292  
vel dapibus odio nulla et ex. Sed viverra cursus mattis. 293  
Suspendisse ornare semper condimentum. Interdum et 294  
malesuada fames ac ante ipsum. 295

**[Paper Contributions]** What is the original con- 296  
tribution of this work? Two or three thoughts that one 297  
should definitely take home. Lorem ipsum dolor sit 298  
amet, consectetur adipiscing elit. Praesent posuere 299  
mattis ante at imperdiet. Cras id tincidunt purus. Ali- 300  
quam erat volutpat. Morbi non gravida nisi, non iaculis 301  
tortor. Quisque at fringilla neque. 302

**[Future Work]** How can other researchers / devel- 303  
opers make use of the results of this work? Do you 304  
have further plans with this work? Or anybody else? 305  
Lorem ipsum dolor sit amet, consectetur adipiscing 306  
elit. Suspendisse sollicitudin posuere massa, non con- 307  
vallis purus ultricies sit amet. Duis at nisl tincidunt, 308  
maximus risus a, aliquet massa. Vestibulum libero 309  
odio, condimentum ut ex non, eleifend. 310

## Acknowledgements 311

I would like to thank my supervisor X. Y. for her help. 312